

Distributed Systems

Slides 36–44 — Complete Academic Summary

◆ Slide 36 — Enterprise Application Integration (EAI)

What is EAI?

Enterprise Application Integration (EAI) is the process of enabling **different enterprise applications** (often developed independently) to **work together** and share data and processes in a coordinated way.

Transactions in EAI

A **transaction** is a sequence of operations that must be executed as a **single logical unit**.

Transaction Primitives

Primitive	Description
BEGIN_TRANSACTION	Marks the start of a transaction
END_TRANSACTION	Terminates the transaction and attempts to commit
ABORT_TRANSACTION	Cancels the transaction and restores old values
READ	Reads data from a file, table, or resource
WRITE	Writes data to a file, table, or resource

The Core Issue: *All-or-Nothing*

- Either all operations succeed, or none of them take effect
 - Partial execution is **not allowed**
-

ACID Properties (Very Important for Exams ⭐)

Atomicity

- The transaction happens **indivisibly**

- Either fully executed or fully aborted

2 Consistency

- The transaction does **not violate system rules or constraints**

3 Isolation

- Concurrent transactions do **not interfere** with each other

4 Durability

- Once committed, changes are **permanent**, even after failures

Exam tip:

ACID guarantees correctness and reliability in distributed enterprise systems.

◆ Slide 37 — TPM: Transaction Processing Monitor

Observation

In many enterprise systems:

- Data involved in a transaction is **distributed across multiple servers**
 - Coordinating these servers is **complex**
-

What is a TPM?

A **Transaction Processing Monitor (TPM)** is middleware responsible for:

- Coordinating the execution of a transaction
 - Managing communication between clients and servers
 - Ensuring ACID properties across distributed resources
-

How TPM Works (Conceptually)

1. Client sends a **transaction request**
2. TPM:
 - Breaks it into sub-requests
 - Sends them to different servers
3. Servers process requests and send replies
4. TPM:
 - Collects responses
 - Decides whether to **commit or abort**
5. Client receives the final result

Exam idea:

TPM hides the complexity of distributed transactions from applications.

◆ Slide 38 — Middleware and EAI

Role of Middleware

Middleware provides **communication facilities** that allow distributed applications to interact and integrate.

It sits **between clients and servers**, abstracting networking and coordination details.

Middleware Communication Models

1 Remote Procedure Call (RPC)

- A client calls a remote operation as if it were local
- The call is:
 - Packaged as a message
 - Sent to the server
 - Processed
 - Result returned as a normal function return

📌 Limitation:

- Caller and callee must be **active at the same time**

2 Message-Oriented Middleware (MOM)

- Messages are:
 - Sent to a **logical contact point** (published)
 - Forwarded to **subscribed applications**
 - Communication is **asynchronous**
- 📌 Advantage:
- Sender and receiver are **decoupled in time and space**

Exam Comparison

RPC → synchronous, tightly coupled
MOM → asynchronous, loosely coupled

- ◆ Slide 39 — How to Integrate Applications

1 File Transfer

- Technically simple
- Not flexible

Problems:

- Need agreement on file format and layout
 - File management complexity
 - Update propagation and notifications are difficult
-

2 Shared Database

- More flexible than file transfer
- Applications share the same data store

Issues:

- Requires a **common data schema**
 - Risk of:
 - Bottlenecks
 - Tight coupling
-

3 Remote Procedure Call (RPC)

- Effective when a **sequence of actions** must be executed
 - Still synchronous
-

4 Messaging

- Allows **decoupling**
- No need for both applications to be running simultaneously
- Better suited for large distributed systems

Exam tip:

Messaging is preferred in scalable, loosely coupled architectures.

- ◆ Slide 40 — Distributed Pervasive Systems

Observation

A new generation of distributed systems where:

- Nodes are **small and mobile**
- Often **embedded** in larger systems
- Systems **blend naturally into the user's environment**

Users may not even be aware they are interacting with a distributed system.

Three Overlapping Subtypes

1 Ubiquitous Computing Systems

- Continuously present
- Continuous interaction between user and system

2 Mobile Computing Systems

- Devices are inherently **mobile**
- Location changes are essential

3 Sensor (and Actuator) Networks

- Emphasis on **collaborative sensing**
- May also perform actions (actuation) on the environment

◆ Slide 41 — Ubiquitous Systems

Core Elements of Ubiquitous Computing

1 Distribution

- Devices are networked, distributed, and transparently accessible

2 Interaction

- User interaction is **highly unobtrusive**
- Technology "fades into the background"

3 Context Awareness

- System understands user context (location, activity, time)
- Uses it to optimize interaction

4 Autonomy

- Devices operate **without human intervention**
- Self-managed systems

5 Intelligence

- System handles complex, dynamic interactions
- Adapts to changing conditions

❖ Key phrase to remember:

Ubiquitous computing is invisible, adaptive, and context-aware.

◆ Slide 42 — Mobile Computing

Distinctive Features

1 Diverse Devices

- Smartphones
 - Tablets
 - GPS devices
 - Remote controls
 - Active badges
-

2 Mobility

- Device location changes over time
- Leads to:
 - Changing local services
 - Dynamic reachability

Keyword: **Discovery**

3 Unstable Communication

- No fixed route
- No guaranteed connectivity
- Leads to **disruption-tolerant networking**

📌 Exam note:

Mobile computing systems must handle intermittent connectivity.

◆ Slide 43 — Sensor Networks

Characteristics of Sensor Nodes

Sensor nodes are:

- **Many**
 - Typically tens to thousands

- **Simple**
 - Limited memory
 - Limited processing power
 - Limited communication capability
- **Energy-constrained**
 - Often battery-powered
 - Sometimes battery-less

❖ **Key challenge:**

Energy efficiency is the dominant design concern in sensor networks.

◆ **Slide 44 — Sensor Networks as Distributed Databases**

Two Extreme Approaches

1 Centralized Approach

- Sensors send **all data directly** to the operator's site
 - Operator stores and processes data
- ✓ Simple
✗ High communication cost
✗ Energy inefficient

2 In-Network Processing Approach

- Each sensor:
 - Processes and stores data locally
 - Operator sends **queries**
 - Sensors return **only answers**, not raw data
- ✓ Energy efficient
✓ Scalable
✓ Reduces communication

❖ **Exam takeaway:**

Treating sensor networks as distributed databases enables efficient querying and aggregation.